
MLP and CNN

Yiannos Georgiou
11807288
University of Amsterdam
ygeorgiou12@gmail.com

1 MLP backprop and NumPy implementation

1.1 Analytical derivation of gradients

1.1.1 Question 1.1.a

1) Cross-Entropy Loss

We can remove the sum because our result is one-hot vector.

$$\begin{aligned}\frac{\partial L}{\partial x^{(N)}} &\Rightarrow \frac{\partial(-\log x_{\text{argmax}(t)}^{(N)})}{\partial x^{(N)}} \\ &\Rightarrow \frac{-1}{x_{\text{argmax}(t)}^{(N)}} \\ \delta_{out} \text{dimensions} &= [1, N]\end{aligned}$$

2) Soft-max Function

Since soft-max is a $N \rightarrow N$ function, the most general derivative we compute for it is the Jacobian matrix.

$$\begin{aligned}\frac{\partial x_i^{(N)}}{\partial \hat{x}_j^{(N)}} &\Rightarrow \frac{\partial(\frac{\exp(\hat{x}^{(N)})}{\sum_{k=1}^{dN} \exp(\hat{x}^{(N)})})}{\partial \hat{x}^{(N)}} \\ &\Rightarrow \frac{(\exp(\hat{x}_i^{(N)}))' * \sum_{k=1}^{dN} \exp(\hat{x}^{(N)}) - \exp(\hat{x}_i^{(N)}) * (\sum_{k=1}^{dN} \exp(\hat{x}^{(N)}))'}{\sum_{k=1}^{dN} \exp(\hat{x}^{(N)})^2} \\ &\Rightarrow \frac{(\exp(\hat{x}_i^{(N)}))' * \sum \hat{x}^{(N)} - \exp(x_i) * \exp(x_j)}{\sum_{k=1}^{dN} \exp(\hat{x}^{(N)})^2} \\ \text{if } i=j &\Rightarrow \frac{\exp(\hat{x}_i^{(N)})}{\sum_{k=1}^{dN} \exp(\hat{x}^{(N)})} \frac{\sum_{k=1}^{dN} \exp(\hat{x}^{(N)}) - \exp \hat{x}_j}{\sum_{k=1}^{dN} \exp(\hat{x}^{(N)})} \\ &\Rightarrow \text{Softmax}(\hat{x}_i)(1 - \text{Softmax}(\hat{x}_j)) \\ \text{if } i \neq j &\Rightarrow \frac{(\exp(\hat{x}_i^{(N)}))' * \sum \hat{x}^{(N)} - \exp(x_i) * \exp(x_j)}{\sum_{k=1}^{dN} \exp(\hat{x}^{(N)})^2} \\ &\Rightarrow \frac{0 * \exp(x_j)}{\sum_{k=1}^{dN} \exp(\hat{x}^{(N)})^2} \\ &\Rightarrow -\text{Softmax}(\hat{x}_i^{(N)}) * \text{Softmax}(\hat{x}_j^{(N)})\end{aligned}$$

The jacobian matrix's dimensions are [N x N] where N is the number of outputs.

10 3) Relu Function

11 For the Relu function we do not need to create a jacobian matrix due to the fact that the output
 12 depends only in a single input.

$$\frac{\partial x_i^{(l < N)}}{\partial x_j^{(l < N)}} \Rightarrow \frac{\partial(\max(0, x_i))}{\partial x_j^{(l < N)}}$$

$$\text{if } i=j \Rightarrow x' \Rightarrow 1$$

$$\text{if } i \neq j \Rightarrow 0 \Rightarrow 0$$

13 4) Linear Function w.r.t. input

$$\frac{\partial \hat{x}_i^{(l)}}{\partial x_j^{(l-1)}} \Rightarrow \frac{\partial(W^{(l)} * x^{(l-1)} + b^{(l)})}{\partial x_j^{(l-1)}}$$

$$\Rightarrow W_j^{(l)}$$

14 5) Linear Function w.r.t. weights

$$\frac{\partial \hat{x}_i^{(l)}}{\partial W_{jk}^{(l)}} \Rightarrow \frac{\partial(W^{(l)} * x^{(l-1)} + b^{(l)})}{\partial W_{jk}^{(l)}}$$

$$\Rightarrow x_k^{(l-1)}$$

15 6) Linear Function w.r.t. biases

$$\frac{\partial \hat{x}_i^{(l)}}{\partial b_j^{(l)}} \Rightarrow \frac{\partial(W^{(l)} * x^{(l-1)} + b^{(l)})}{\partial b_j^{(l)}}$$

$$\text{if } i=j \Rightarrow b_{ij}' \Rightarrow 1$$

$$\text{if } i \neq j \Rightarrow 0$$

16 **1.1.2 Question 1.1.b**

17 1)

$$\frac{\partial L}{\partial \hat{x}^{(N)}} = \frac{\partial L}{\partial x^{(N)}} \frac{\partial x^{(N)}}{\partial \hat{x}^{(N)}} \Rightarrow \delta_{out}^{(N)} \cdot \text{Jacobian_matrix}$$

$$\Rightarrow \delta_{N \times 1}^{(N)}$$

$$\text{Jacobian_matrix} = \begin{bmatrix} S_1(\delta_{11} - S_1) & S_1(\delta_{12} - S_2) & S_1(\delta_{13} - S_3) & \dots & S_1(\delta_{1N} - S_N) \\ \dots & \dots & \dots & \dots & \dots \\ S_N(\delta_{N1} - S_1) & S_N(\delta_{N2} - S_2) & S_N(\delta_{N3} - S_3) & \dots & S_N(\delta_{NN} - S_N) \end{bmatrix}$$

$$\text{if } i=j \delta_{ij} = 1$$

$$\text{else } \delta_{ij} = 0$$

18 2)

$$\frac{\partial L}{\partial \hat{x}^{l < N}} = \frac{\partial L}{\partial x^{(l)}} \frac{\partial x^{(l)}}{\partial \hat{x}^{(l)}} \Rightarrow \delta_{out}^{(l)} * \text{relu}'(x)$$

$$\frac{\partial L}{\partial \hat{x}^{(l)}} \Rightarrow \begin{bmatrix} \delta_1 * \text{relu}'(x_1) \\ \delta_2 * \text{relu}'(x_2) \\ \dots \\ \delta_{dim_l} * \text{relu}'(x_l) \end{bmatrix}$$

19 3) Linear backward w.r.t inputs

$$\begin{aligned}\frac{\partial L}{\partial x^{(l+1)}} &= \frac{\partial L}{\partial \hat{x}^{(l+1)}} \frac{\partial \hat{x}^{(l+1)}}{\partial x^{(l)}} \Rightarrow \delta_{dim_{l+1} \times 1}^{(l+1)} \cdot \frac{\partial (W^{(l+1)}x^{(l)} + b^{l+1})}{\partial x^{(l)}} \\ &\Rightarrow W_{d_{l+1} \times d_l, d_{l+1} \times 1}^T \delta_{d_{l+1} \times 1}^{(l+1)} \Rightarrow \delta_{d_l \times 1}^l\end{aligned}$$

20 4) Linear backward w.r.t. weights

$$\begin{aligned}\frac{\partial L}{\partial W^l} &= \frac{\partial L}{\partial \hat{x}^{(l)}} \frac{\partial \hat{x}^{(l)}}{\partial W^{(l)}} \Rightarrow \delta^{(l+1)} \cdot \frac{\partial (W^{(l)}x^{(l-1)} + b^{(l)})}{\partial W^{(l)}} \\ &\Rightarrow \delta_{d_l \times 1}^l \cdot x_{d_l \times 1}^l \Rightarrow x_{d_l \times 1}^l \cdot (\delta^l)^T \Rightarrow \frac{\partial L}{\partial W_{d_l \times d_{l+1}}}\end{aligned}$$

21 5) Linear backward w.r.t. biases

$$\begin{aligned}\frac{\partial L}{\partial b^{(l)}} &= \frac{\partial L}{\partial \hat{x}^{(l)}} \frac{\partial \hat{x}^{(l)}}{\partial b^{(l)}} \Rightarrow \delta^{(l)} \cdot \frac{\partial (W^{(l)}x^{(l-1)} + b^{(l)})}{\partial b^{(l)}} \\ &\Rightarrow \delta_{d_l \times 1}^{(l)} \cdot 1 \Rightarrow \frac{\partial L}{\partial b_{d_l \times 1}^{(l)}}\end{aligned}$$

22 1.1.3 Question 1.1.c

23 The difference is that the loss derivatives between the layers are matrices if $B > 1$, so some equations
24 have to be changed and summed up to fit on the back-propagation or on the updates of the weights.
25 For instance, on the update of the biases you have to sum the batch to get a vector, or in cases such as
26 soft-max where if your batch size is more than 1 you have to multiply with a 3d-array. However, batch
27 size it doesn't make a lot of difference in learning or testing procedures, except in the implementation.

28 1.2 NumPy implementation

29 NumPy implementation based on the equations derived from the sections 1.1.a and 1.1.b. Test
30 accuracy reached 46.5% with default parameter setup. Test and train accuracies can be seen in Figure
31 1.

32 2 PyTorch MLP

33 Experiments were focused on two optimizers, ADAM and SGD. illustrated on Figure 2 are the best
34 models for these optimization algorithms. As it can be seen the ADAM optimizer outperform SGD.
35 However, different parameters and architectures boost the performance of each one.

36 The architecture of our best SGD model contained 3 hidden layers with 300 neuron each, the learning
37 rate was initialized to 0.05 and was reducing to half every 2000 iterations.

38 Regarding ADAM best model, its architecture contained 2 hidden layers with 300 and 200 neurons
39 respectively. The learning rate was initialized to 0.0006 and was reducing to half every 2000 iterations.
40 Batch normalization seem benefit this model.

41 **Batch normalization** : Batch normalization works well only for ADAM, and boost its accuracy to
42 51.5%. Although, that does not apply to SGD where the accuracy after the batch norm is getting
43 worst.

44 **Batch size**: Although we kept the original batch size, we experimented with higher and lower values
45 for this parameter. We observed that high batch size cause a drop of accuracy while it makes the
46 model faster, and vice versa.



Figure 1: MLP loss and accuracy measurements (train accuracy (up,left), train loss (up,right), test accuracy(down,middle))

47 **Learning rate:** The current configuration is not the ideal for our model, so we tried to find the
 48 best learning rate for each model. The learning rate is affected by the size of the network and the
 49 optimization algorithm. We noticed that, if the network is too big is preferable to have a small
 50 learning rate. Moreover, higher accuracy was achieved when the learning rate was reducing over the
 51 iterations.

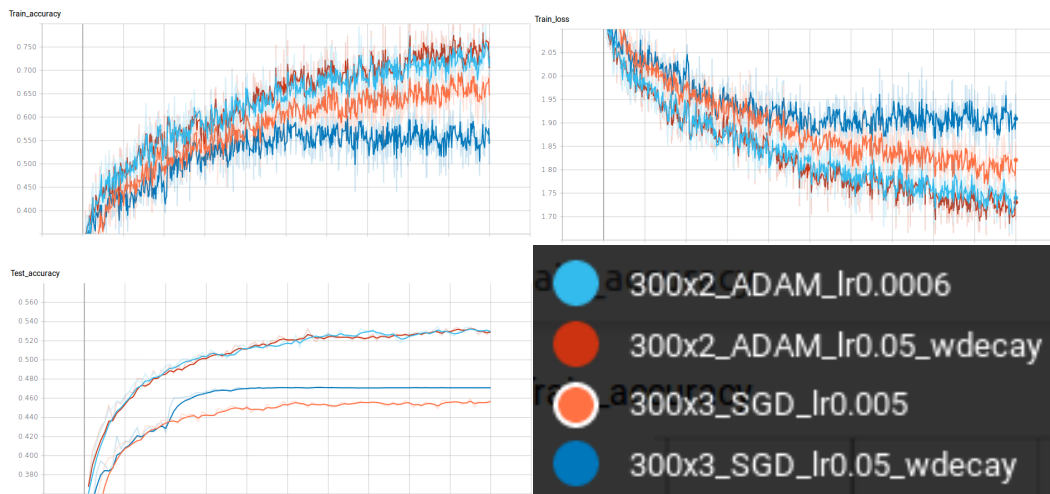


Figure 2: MLP pytorch loss and accuracy measurements (train accuracy (up,left), train loss (up,right), test accuracy(down,middle))

52 3 Custom Module: Batch Normalization

53 3.1 Automatic differentiation

54 3.2 Manual implementation of backward pass

55 1) gamma derivative

$$\begin{aligned}\frac{\partial L}{\partial \gamma} &= \frac{\partial L}{\partial y} \frac{\partial y}{\partial \gamma} \Rightarrow \delta_{out} * \frac{\partial(\gamma * \hat{x} + \beta)}{\partial \gamma} \\ &\Rightarrow \delta_{out} * \frac{\hat{x}}{N \times D} \Rightarrow \sum_{n=1}^N \delta_{out n} * \hat{x}_n \Rightarrow \frac{\partial L}{\frac{\partial \gamma}{D \times 1}}\end{aligned}$$

56 2) beta derivative

$$\frac{\partial L}{\partial \beta} = \frac{\partial L}{\partial y} \frac{\partial y}{\partial \beta} \Rightarrow \delta_{out} * \frac{\partial(\gamma * \hat{x} + \beta)}{\partial \beta} \Rightarrow \sum_{n=1}^N \delta_{out n} * 1$$

57 3) x derivative

$$\begin{aligned}\frac{\partial L}{\partial x} &= \frac{\partial L}{\partial y} \frac{\partial y}{\partial x} \Rightarrow \delta_{out} * \frac{\partial(\gamma * \hat{x} + \beta)}{\partial x} \Rightarrow \delta_{out} * \frac{\partial(\gamma * \hat{x} + \beta)}{\partial \hat{x}} \frac{\partial \hat{x}}{x} \Rightarrow \delta_{out} * \frac{\partial(\gamma * \hat{x} + \beta)}{\partial \hat{x}} \frac{\partial(x - mean) * \frac{1}{\sqrt{\sigma^2 + \epsilon}}}{\partial(x - mean)} \\ &\Rightarrow dxm_1 = \delta_{out} * \gamma \frac{1}{\sqrt{\sigma^2 + \epsilon}} \\ &\Rightarrow \delta_{out} * \gamma \frac{\partial(x - mean) * \frac{1}{\sqrt{\sigma^2 + \epsilon}}}{(\frac{1}{\sqrt{\sigma^2 + \epsilon}})} \Rightarrow \delta_{out} * \gamma * (x - mean) * \frac{\partial(\frac{1}{\sqrt{\sigma^2 + \epsilon}})}{\partial \sqrt{\sigma^2 + \epsilon}} \\ &\Rightarrow \delta_{out} * \gamma * (x - mean) * \frac{1}{\sqrt{\sigma^2 + \epsilon}^2} * \frac{\partial \sqrt{\sigma^2 + \epsilon}}{\partial \sigma^2} \\ &\Rightarrow \delta_{out} * \gamma * (x - mean) * \frac{1}{\sqrt{\sigma^2 + \epsilon}^2} * 0.5 * \frac{1}{\sqrt{\sigma^2 + \epsilon}} * \frac{\partial \frac{1}{B} \sum_{s=1}^B (x - mean)^2}{\partial(x - mean)} \\ &\Rightarrow dxm_2 = \delta_{out} * \gamma * (x - mean) * \frac{1}{\sqrt{\sigma^2 + \epsilon}^2} * 0.5 * \frac{1}{\sqrt{\sigma^2 + \epsilon}} * \frac{2}{B} * (x - mean) * \frac{\partial(x - mean)}{\partial x} \\ dx &= dxm_1 + dxm_2\end{aligned}$$

58 4 PyTorch CNN

59 Convolutional networks proved to be undoubtedly more suitable for this task by observing the derived
60 results in Figure 3. The thing is that convolution networks have the ability to exploit spatially local
61 correlation present in natural images. We can observe that convolution network outperforms MLP in
62 both accuracy and loss.

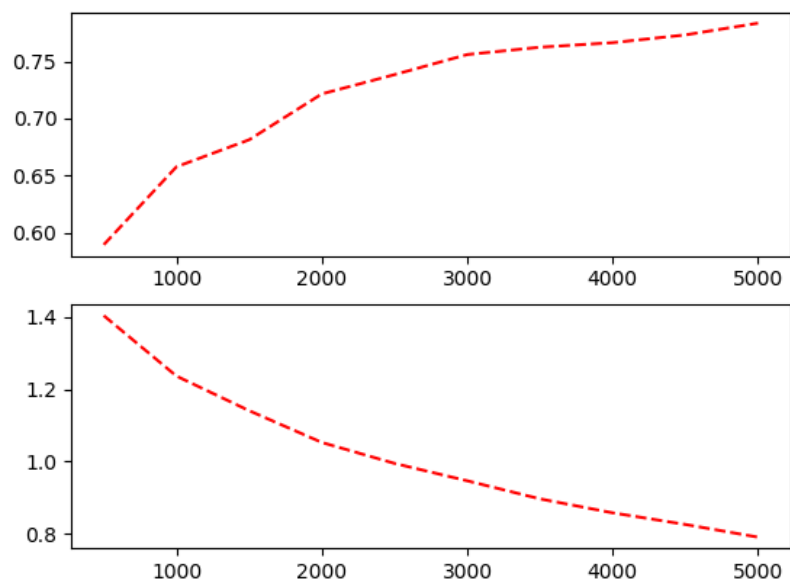


Figure 3: Convolution network loss and accuracy measurements (accuracy (up), loss (down))